Module 1 knowledge check

A developer is designing a container for an AI application that stores user interaction logs. Each document includes a userId property. The application frequently retrieves all logs for a specific user. Which partition key provides the best performance?

A. Use userId as the partition key
B. Use a timestamp property as the partition key
C. Use a boolean isProcessed property as the partition key

An AI application stores product recommendations with document IDs. The application needs to retrieve a specific recommendation by its known ID and category. Which method provides the most efficient retrieval?

A. Use query_items() with a WHERE clause filtering by ID
B. Use read_item() with the item ID and partition key
C. Use query_items() with enable_cross_partition_query=True

Module 2 knowledge check

A developer creates a container for documents with embeddings from the text-embedding-ada-002 model. Which vector policy configuration is correct for this model?

A. Set dataType to float32, dimensions to 1536, and distanceFunction to cosine
B. Set dataType to int8, dimensions to 1024, and distanceFunction to euclidean
C. Set dataType to float16, dimensions to 3072, and distanceFunction to dotproduct

A team wants search that combines semantic understanding with exact keyword matching for error codes. Which approach enables hybrid search?

A. Use ORDER BY RANK RRF with VectorDistance and FullTextScore functions
B. Execute two separate queries and merge results in application code
C. Add keyword terms to the query vector before executing vector search

Module 3 knowledge check

A developer needs to optimize a query that filters by documentType and sorts by uploadDate descending. Which index configuration best supports this pattern?

A. Create a composite index with documentType (ascending) followed by uploadDate (descending)
B. Create separate range indexes on documentType and uploadDate
C. Create a composite index with uploadDate (descending) followed by documentType (ascending)

Query metrics show low index utilization and high retrieved-to-output document ratios. What does this indicate?

A. The query is performing scans instead of using indexes. Add range or composite indexes for the query's filter and sort properties.
B. The container has too many indexes. Remove indexes to reduce query overhead.
C. The query returns too many results. Add a TOP clause to limit output